

# VAE and GRU Assignment

## Assignment: Time Series Forecasting and Anomaly Detection with Weather Data

In this assignment, you'll implement both a Variational Autoencoder (VAE) and a Gated Recurrent Unit (GRU) network to analyze weather data. You'll use the VAE for anomaly detection and the GRU for time series forecasting.

### Dataset

You'll work with the "Hourly Weather Surface - Brazil" dataset, which contains hourly weather measurements from Brazil including temperature, humidity, wind speed, and other atmospheric conditions.

### Part 1: Data Preparation

Download the dataset from:

<https://www.kaggle.com/datasets/PROPPG-PPG/hourly-weather-surface-brazil-southeast-region>

1. Perform exploratory data analysis:
  - Visualize distributions of key features
  - Create time series plots for temperature, humidity, and wind speed
  - Check for missing values and outliers
2. Preprocess the data:
  - Handle missing values
  - Normalize/standardize features
  - Create sequences for time series forecasting
  - Split into training, validation, and test sets

### Part 2: VAE Implementation for Anomaly Detection

1. Implement a VAE architecture with:
  - An encoder with at least 2 hidden layers
  - A latent space dimension of your choice (justify your selection)
  - A decoder with at least 2 hidden layers
  - Proper sampling from the latent distribution
2. Train the VAE on normal weather patterns:
  - Use appropriate loss functions (reconstruction loss + KL divergence)
  - Implement early stopping
  - Visualize the training and validation loss curves
3. Anomaly detection:
  - Define an anomaly score based on reconstruction error
  - Detect anomalous weather days (e.g., extreme temperature changes)

- Visualize the anomalies on a time series plot
  - Evaluate your anomaly detection performance using precision, recall, and F1-score
  - Consider different threshold values for anomaly detection and justify your choice
4. Latent space analysis:
- Visualize the latent space (use dimensionality reduction if needed)
  - Analyze how different weather patterns are represented in the latent space
  - Generate new synthetic weather data by sampling from the latent space

### **Part 3: GRU Implementation for Time Series Forecasting**

1. Implement a GRU-based network:
  - Input features should include multiple weather metrics
  - At least one GRU layer with justified number of units
  - Appropriate output layer for forecasting
2. Configure the model for sequence-to-sequence forecasting:
  - Use past N hours to predict the next M hours (choose appropriate N and M)
  - Implement teacher forcing during training (with a decay schedule)
3. Train the GRU model:
  - Implement batch training
  - Use an appropriate loss function for regression
  - Implement learning rate scheduling
  - Visualize training and validation loss
4. Evaluate forecasting performance:
  - Compute MSE, MAE, and RMSE on the test set
  - Visualize predictions vs. actual values
    - Analyze how far into the future the model can predict accurately
5. Ablation study:
  - Compare GRU performance to a simpler LSTM and RNN
  - Analyze how changing the sequence length affects performance
  - Investigate how different features impact prediction accuracy

### **Part 4: Model Integration and Analysis**

1. Integrate the VAE and GRU models:
  - Use the VAE to detect anomalies in the input data
  - Analyze how the GRU model performs on normal vs. anomalous data